

Rejection sampling

Computational Statistics

University of Copenhagen

September 25th 2019

Introduction

Sample from the probability distribution on $[0; \infty[$ with a density

$$f(y) \propto \prod_{i=1}^{100} \exp(yz_i x_i - e^{yx_i}), \quad y \geq 0$$

Find a Gaussian envelope of f and implement rejection sampling from the distribution with density f using this envelope.

Implement the adaptive rejection sampling algorithm that uses a piecewise log-affine envelope and compare it with the one based on the Gaussian envelope

Target density and implementation

$$f(y) \propto \prod_{i=1}^{100} \exp(yz_i x_i - e^{yx_i}), \quad y \geq 0$$


```
target_sapply <- function(y)
{
  sapply(y, function(yy) prod(exp(yy*zx - exp(x*yy)))) / 1.05009e-41
}

target_sum <- function(y)
{
  sapply(y, function(yy) exp( sum(yy*zx - exp(x*yy)) )) / 1.05009e-41
}


target_outer <- function(y)
{
  exp( rowSums(outer(y, zx) - exp(outer(y, x)))) / 1.05009e-41
}
```

Note that the normalization constant is approximated via *integrate()*, so the α -notation still applies. It does however allow us to later interpret α as being approximately the true probability of accept.

Comparisons




Target density




Gaussian envelope

```
mean <- integrate(function(v) target_outer(v)*v, 0, Inf)$value  
sd <- sqrt(integrate(function(v) target_outer(v)*v^2, 0, Inf)$value - mean^2)
```




Ratio plot

The α -parameter was calculated to be approximately 0.237.




The envelope




Improved envelope

Clearly, our proposal needs to have a higher standard deviation. By trial and error, we increase the sd by 9 percent.




Improved envelope

The new α was calculated to be 0.9163.




Improved envelope



A (generalized) t-distributed envelope

The α here is 0.954. $\mu = 0.24$, $\sigma = 0.058$, $df = 61$.

We found it by simulating the density from the Gaussian envelope, then fitting a t-distribution on the results.



Rejection sampling algorithm slow

#We encapsulate all objects needed in a list.

#We will re-use this format later for the adaptive envelope generator

#this allows for easy extension to OO-programming

```
proposal_factory <- function(target_dens, proposal_dens, proposal_sim, a)
{
  list( tdens = target_dens,
        pdens = proposal_dens,
        sim = proposal_sim,
        alpha = a)}

target_simulator_slow <- function(n, proposal_object) {
  y <- numeric(n)
  for(i in 1:n) {
    reject <- TRUE
    while(reject) {
      y0 <- proposal_object$sim(1)
      u <- runif(1)
      reject <- u > proposal_object$a*proposal_object$tdens(y0) /
        proposal_object$pdens(y0)
    }
    y[i] <- y0
  }
  y}
```

Rejection sampling algorithm improved

```
target_simulator <- function(n, proposal_object, scale = 1 )
{
  simulated_values <- numeric(n) #we need n samples

  num_accepted <- 0 #we keep track of how many we have

  while (num_accepted < n)
  {
    samples <- ceiling( scale*(n - num_accepted)/proposal_object$a )
    #sample expected amount but possibly scaled up by a factor

    prop_samples <- proposal_object$sim(samples) #get all the samples
    uniform_samples <- runif(samples)

    accept <- uniform_samples <= proposal_object$a*proposal_object$tdens(prop_samples) /
      proposal_object$pdens(prop_samples)


    new_additions <- min(n - num_accepted, sum(accept))

    simulated_values[(num_accepted + 1):(num_accepted + new_additions)]
      <- (prop_samples[accept])[1:new_additions]


    num_accepted <- num_accepted + new_additions
  }

  simulated_values
}
```

Target simulator slow




Target simulator improved



Comparing the Gaussian and t -distributed envelopes

Unit: milliseconds

	expr	min	lq	mean	median	uq
target_simulator(10000, propn, 1)		165.3658	167.6069	180.1751	172.6749	187.1796
target_simulator(10000, propt, 1)		164.4175	165.5598	176.6923	171.1025	180.7648
max neval						
273.6954		100				
268.6328		100				




The benefit of a higher accept probability is largely cancelled out by the slower simulation and evaluation of the t -envelope.

Main implementation with parallelization


A makeshift example, using the doParallel library:

```
foreach(i = rep(10000/8, 8), .combine = 'c') %dopar%  
{  
  propn <- proposal_factory(target_outer, function(x) dnorm(x, mean, sd*1.09),  
    function(x) rnorm(x, mean, sd*1.09), alpha)  
  target_simulator(i, propn, 1)  
}
```

Comparing the implementations, $n = 10.000$



Comparing the implementations, $n = 50.000$



Can Rcpp speed things up?

```
// [[Rcpp::export]]
NumericVector target_simulator_cpp(int n, List prop_object, double scale)
{
  Function target = as<Function>(prop_object["tdens"]);
  Function proposal = as<Function>(prop_object["pdens"]);
  Function sim = as<Function>(prop_object["sim"]);
  double alpha = as<double>(prop_object["alpha"]);

  NumericVector simulated_values(n);

  int num_accepted = 0;
  while (num_accepted < n)
  {
    int samples = ceil(scale*(double)(n - num_accepted)) / alpha;
    NumericVector prop_samples = sim(samples);
    NumericVector uniform_samples = runif(samples);


    NumericVector fx = target(prop_samples);
    NumericVector gx = proposal(prop_samples);

    for (int i = 0; i < samples; i++)
    {
      if (alpha*fx[i] / gx[i] >= uniform_samples[i])
      {
        simulated_values[num_accepted] = prop_samples[i];
        num_accepted += 1;
        if (num_accepted >= n)
        {
          goto exit;
        }
      }
    }
  }
exit:
  return simulated_values;
}
```

Unit: milliseconds

	expr	min	lq	mean	median	uq	max	neval
target_simulator(5000, propn, 1)		82.96958	83.35984	87.55625	83.86785	93.26021	99.72847	100
target_simulator_cpp(5000, propn, 1)		82.67603	83.09416	89.41752	83.52766	89.09366	222.22728	100

Comparing the main implementation for different values of *scale*, $n = 10000$



Profiling ($n = 100.000$, Gaussian envelope)

```
profvis({
  simulated_values <- numeric(n) #we need n samples

  num_accepted <- 0 #we keep track of how many we have

  while (num_accepted < n)
  {
    samples <- ceiling( scale*(n - num_accepted)/proposal_object$a ) #sample expected amount

    prop_samples <- proposal_object$sim(samples) #get all the samples
    uniform_samples <- runif(samples)

    accept <- uniform_samples <= proposal_object$a*proposal_object$tdens(prop_samples) /
      proposal_object$pdens(prop_samples)

    new_additions <- min(n - num_accepted, sum(accept))

    simulated_values[(num_accepted + 1):(num_accepted + new_additions)] <- (prop_samples[accept])
    [1:new_additions]

    num_accepted <- num_accepted + new_additions
  }

  simulated_values()
}
```

0.1	20
0.9	20
-1.1	169.0 1740
3.1	10

Vast majority of time is spent in the evaluation of the densities. It can further be shown that it is the target density which is causing the problem.

This specific density could be implemented in C++, but the generic parts of the code seem otherwise well-optimized.

Adaptive envelope factory

```
envelope_factory <- function(x,
                             target_dens,
                             logderiv = NULL,
                             lower_support = -Inf,
                             upper_support = Inf)
{
  #check if the log-derivative was supplied, otherwise we get it ourselves:
  if (is.null(logderiv))
  {
    logderiv <- function(xx) grad(function(x) log(target_dens(x)), xx)
  }

  #calculate a-vector and stop if there are integrability issues

  a <- logderiv(x)

  #check if we need to stop
  continue <- (a[1] > 0 & a[length(a)] < 0) | (a[length(a)] > 0 & lower_support > -Inf)
  | a[1] < 0 & upper_support < Inf

  if (!continue)
  {
    stop("Envelope is not integrable. Re-submit new x")
  }

  #now calculate b, z, Fz, Q, and const

  b <- log(target_dens(x)) - a*x

  z <- c(lower_support, -diff(b)/diff(a), upper_support)

  Fz <- numeric(length(x))

  for (i in seq_along(Fz))
  {
    Fz[i] <- exp(b[i]) * ( exp( a[i] * z[i+1] ) - exp(a[i]*z[i])) / a[i]
  }

  Q <- c(0, cumsum(Fz))
}
```

Adaptive envelope factory continued

```
const <- Q[length(Q)]
```

```
#now define simulator and envelope -- must be vectorized
```

```
proposal_density <- function(x)
```

```
  index <- findInterval(x, z) #Given x, find z-interval that x belongs to
```

```
#findInterval is optimized and  $O(\log(\text{length}(z)*\text{length}(x)))$ 
```

```
#now just evaluate the function and return
```

```
  V <- a[index]*x + b[index]
```

```
  exp(V)/const
```

```
}
```

```
simulator <- function(n)
```

```
{
```

```
  #we need n uniform samples
```

```
  q <- runif(n)
```

```
#find the index to which const*q belongs to.
```

```
  index <- findInterval(const*q, Q) + 1
```

```
#solve the equation for x
```

```
  log( (const*q - Q[index])*a[index ]*exp(-b[index])
```







```
  + exp(a[index]*z[index]))/a[index ]
```

```
}
```





```
#return the list-object that can be used directly in previous functions.
```

```
list( tdens = dens$tdens, pdens = proposal_density, sim = simulator, alpha = 1/const)
```


Comparing envelopes (3 to 8 nodes)



Testing the generator for various lengths of the x-vector




Timing the generator for various lengths of the x-vector




Comparing the adaptive envelope ($\alpha = 0.9839$) with gaussian envelope ($\alpha = 0.9163$), $n = 20000$

Adaptive envelope




Gaussian envelope




Comparing the adaptive envelope ($\alpha = 0.9839$) with
gaussian envelope ($\alpha = 0.9163$), $n = 20000$

Ratio of empirical CDFs (easily passes the Kolmogorov-Smirnov test)



Comparing the adaptive envelope ($\alpha = 0.9839$) with
gaussian envelope ($\alpha = 0.9163$), $n = 10000$



Using OOP

```
proposal_factory <- function(target_dens, proposal_dens, proposal_sim, a)
{
  structure(list( tdens = target_dens,
                  pdens = proposal_dens,
                  sim = proposal_sim,
                  alpha = a), class = "RS")
}

plot.RS <- function(object, support)
{
  gridx <- seq(support[1], support[2], length.out = 512)

  plot(gridx, object$alpha*object$tdens(gridx), col = "#286e99", xlab = "", ylab = "",
        main = paste("Alpha equals", propn$alpha), type = "l", lwd = 2)

  points(gridx, object$pdens(gridx), col = "#a89e82")
}

simulate.RS <- function(object, nsim = 1, scale = 1, plot = FALSE)
{
  simulations <- target_simulator(nsim, object, scale)

  if (plot)
  {
    hist(simulations, prob = TRUE, main = "", ylab = "", xlab = "", col = "#a89e82")
    curve(object$tdens(x), add = TRUE)
  }
  simulations
}
```